

In [1]:

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt

np.random.seed(1234)
```

In [2]:

```
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.fashion_mnist.load_data()
x_val = x_train[50000:60000]
x_train = x_train[0:50000]
y_val = y_train[50000:60000]
y_train = y_train[0:50000]
x_train = x_train.astype(np.float32).reshape(-1,28,28,1) / 255.0
x_val = x_val.astype(np.float32).reshape(-1,28,28,1) / 255.0
x_test = x_test.astype(np.float32).reshape(-1,28,28,1) / 255.0
y_train = tf.one_hot(y_train, depth=10)
y_val = tf.one_hot(y_val, depth=10)
y_test = tf.one_hot(y_test, depth=10)
print(x_train.shape)
print(x_test.shape)
print(x_val.shape)
train_dataset = tf.data.Dataset.from_tensor_slices((x_train, y_train))
train_dataset = train_dataset.shuffle(buffer_size=1024).batch(128)
train_dataset_full = train_dataset.shuffle(buffer_size=1024).batch(len(train_dataset))
val_dataset = tf.data.Dataset.from_tensor_slices((x_val, y_val))
val_dataset = val_dataset.batch(128)
test_dataset = tf.data.Dataset.from_tensor_slices((x_test, y_test))
test_dataset = test_dataset.batch(128)
print(len(train_dataset))
print(len(test_dataset))
```

```
2022-03-22 19:56:37.862990: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
```

```
(50000, 28, 28, 1)
```

```
(10000, 28, 28, 1)
```

```
(10000, 28, 28, 1)
```

```
391
```

```
79
```

```
2022-03-22 19:56:37.872284: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
```

```
2022-03-22 19:56:37.872749: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
```

```
2022-03-22 19:56:37.874761: I tensorflow/core/platform/cpu_feature_guard.cc:151] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA
```

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
2022-03-22 19:56:37.876484: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
```

```
2022-03-22 19:56:37.876931: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
```

```
2022-03-22 19:56:37.877253: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
```

```
2022-03-22 19:56:38.254126: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
```

```
2022-03-22 19:56:38.254419: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
```

```
2022-03-22 19:56:38.254792: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
```

```
2022-03-22 19:56:38.255107: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1525] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 7014 MB memory: -> device: 0, name: NVIDIA GeForce GTX 1070, pci bus id: 0000:01:00.0, compute capability: 6.1
```

In [3]:

```

class BatchNormalization(tf.keras.layers.Layer):
    def __init__(self, batch_size, training=False):
        super(BatchNormalization, self).__init__()

        self.convexCoeff = 0.9
        self.numCalls = 0
        self.batch_size = batch_size
        self.training = training

        self.gamma = self.add_weight(name='gamma', shape=[self.batch_size,], ini-
alizer=tf.initializers.ones, trainable=True)
        self.beta = self.add_weight(name='beta', shape=[self.batch_size,], ini-
alizer=tf.initializers.zeros, trainable=True)
        self.mean = self.add_weight(name='mean', shape=[self.batch_size,], ini-
alizer=tf.initializers.zeros, trainable=False)
        self.var = self.add_weight(name='var', shape=[self.batch_size,], ini-
alizer=tf.initializers.zeros, trainable=False)

    def batch_norm(self, inputs, training):
        self.numCalls += 1

        axes = list(range(len(inputs.shape) - 1))
        mean = tf.reduce_mean(inputs, axes, keepdims=True)
        var = tf.reduce_mean(tf.math.squared_difference(inputs, tf.stop_gradient(
mean)), axes, keepdims=True)

        if training:
            norm = tf.add(tf.multiply(self.gamma, tf.divide(tf.subtract(inputs, me-
an), tf.sqrt(var+1e-7))), self.beta)

            mean = tf.squeeze(mean, axes)
            var = tf.squeeze(var, axes)

            moving_avg_mean = ((self.convexCoeff/self.numCalls)*mean) + (1-(self.c-
onvexCoeff/self.numCalls)*self.mean)
            moving_avg_var = ((self.convexCoeff/self.numCalls)*var) + (1-(self.c-
onvexCoeff/self.numCalls)*self.var)

            self.mean.assign(moving_avg_mean)
            self.var.assign(moving_avg_var)

        else:
            norm = tf.add(tf.multiply(self.gamma, tf.divide(tf.subtract(inputs, me-
an), tf.sqrt(var+1e-7))), self.beta)

        return norm

```

In [4]:

```

class ImageRecognitionCNN(tf.keras.Model):

    def __init__(self, num_classes, device='cpu:0', checkpoint_directory=None):
        """ Define the parameterized layers used during forward-pass, the device
            where you would like to run the computation (GPU, TPU, CPU) on and the
            checkpoint directory.

            Args:
                num_classes: the number of labels in the network.
                device: string, 'cpu:n' or 'gpu:n' (n can vary). Default, 'cpu:0'.
                checkpoint_directory: the directory where you would like to save o

            ...
            restore a model.
        """
        super(ImageRecognitionCNN, self).__init__()

        # Initialize layers
        self.conv1 = tf.keras.layers.Conv2D(64, 3, padding='same', activation=None)
        self.conv2 = tf.keras.layers.Conv2D(64, 3, padding='same', activation=None)
        self.pool1 = tf.keras.layers.MaxPool2D()
        self.conv3 = tf.keras.layers.Conv2D(64, 3, padding='same', activation=None)
        self.conv4 = tf.keras.layers.Conv2D(64, 3, padding='same', activation=None)
        # self.pool2 = tf.keras.layers.MaxPool2D()
        # self.conv5 = tf.keras.layers.Conv2D(64, 3, padding='same', activation=None)
        # self.pool2 = tf.keras.layers.MaxPool2D()
        # self.conv6 = tf.keras.layers.Conv2D(64, 3, 2, padding='same', activation=None)
        # self.conv7 = tf.keras.layers.Conv2D(64, 1, padding='same', activation=None)
        self.conv8 = tf.keras.layers.Conv2D(num_classes, 1, padding='same', activation=None)
        self.BN = BatchNormalization(64)

        # Define the device
        self.device = device

        # Define the checkpoint directory
        self.checkpoint_directory = checkpoint_directory
        self.acc = tf.keras.metrics.Accuracy()

    def predict(self, images, training):
        """ Predicts the probability of each class, based on the input sample.

            Args:
                images: 4D tensor. Either an image or a batch of images.
                training: Boolean. Either the network is predicting in
                    training mode or not.

```

```

"""
x = self.conv1(images)
x = self.BN.batch_norm(x, training)
x = tf.nn.relu(x)
x = self.pool1(x)
x = self.conv2(x)
x = self.BN.batch_norm(x, training)
x = tf.nn.relu(x)
x = self.pool1(x)
x = self.conv3(x)
x = self.BN.batch_norm(x, training)
x = tf.nn.relu(x)
x = self.pool1(x)
x = self.conv4(x)
x = self.BN.batch_norm(x, training)
x = tf.nn.relu(x)
x = self.pool1(x)
x = self.conv8(x)
#x = tf.nn.relu(x)
#print(x.shape)
x = tf.reshape(x, (-1, 1, 10))
#x = tf.keras.layers.Flatten(x)
return x

```

```

def loss_fn(self, images, target, training):
    """ Defines the loss function used during
        training.
    """
    preds = self.predict(images, training)
    #print(preds.shape)
    #print(target.shape)
    loss = tf.nn.softmax_cross_entropy_with_logits(labels=target, logits=preds
)
    return loss

```

```

def grads_fn(self, images, target, training):
    """ Dynamically computes the gradients of the loss value
        with respect to the parameters of the model, in each
        forward pass.
    """
    with tf.GradientTape() as tape:
        loss = self.loss_fn(images, target, training)
    return tape.gradient(loss, self.variables)

```

```

def restore_model(self):
    """ Function to restore trained model.
    """
    with tf.device(self.device):
        # Run the model once to initialize variables
        dummy_input = tf.constant(tf.zeros((1,48,48,1)))
        dummy_pred = self.predict(dummy_input, training=False)
        # Restore the variables of the model
        saver = tf.Saver(self.variables)

```

```

        saver.restore(tf.train.latest_checkpoint
                      (self.checkpoint_directory))

def save_model(self, global_step=0):
    """ Function to save trained model.
    """
    tf.Saver(self.variables).save(self.checkpoint_directory,
                                  global_step=global_step)

# def compute_accuracy(self, input_data):
#     """ Compute the accuracy on the input data.
#     """
#     with tf.device(self.device):
#         #acc = tf.metrics.Accuracy()
#         for step, (images, targets) in enumerate(input_data):
#             # Predict the probability of each class
#             #print(targets.shape)
#             logits = self.predict(images, training=False)
#             # Select the class with the highest probability
#             #print(logits.shape)
#             logits = tf.nn.softmax(logits)
#             logits = tf.reshape(logits, [-1, 10])
#             targets = tf.reshape(targets, [-1, 10])
#             preds = tf.argmax(logits, axis=1)

#             #m1.update_state
#             # Compute the accuracy
#             #print(preds.shape)
#             acc(tf.reshape(targets, preds))
#     return acc

def compute_accuracy_2(self, images, targets, training=False):
    """ Compute the accuracy on the input data.
    """
    with tf.device(self.device):

        # Predict the probability of each class
        logits = self.predict(images, training)
        # Select the class with the highest probability

        logits = tf.nn.softmax(logits)
        logits = tf.reshape(logits, [-1, 10])
        targets = tf.reshape(targets, [-1, 10])
        preds = tf.argmax(logits, axis=1)
        goal = tf.argmax(targets, axis=1)
        self.acc.update_state(goal, preds)
        # Compute the accuracy
        result = self.acc.result().numpy()
    return result

def fit_fc(self, training_data, eval_data, optimizer, num_epochs=500,
           early_stopping_rounds=10, verbose=10, train_from_scratch=False):
    """ Function to train the model, using the selected optimizer and
    for the desired number of epochs. You can either train from scratch
    or load the latest model trained. Early stopping is used in order to

```

*mitigate the risk of overfitting the network.*

**Args:**

*training\_data: the data you would like to train the model on.  
Must be in the tf.data.Dataset format.*

*eval\_data: the data you would like to evaluate the model on.  
Must be in the tf.data.Dataset format.*

*optimizer: the optimizer used during training.*

*num\_epochs: the maximum number of iterations you would like to  
train the model.*

*early\_stopping\_rounds: stop training if the loss on the eval  
dataset does not decrease after n epochs.*

*verbose: int. Specify how often to print the loss value of the net*

*work.*

*train\_from\_scratch: boolean. Whether to initialize variables of th*

*e*  
*the last trained model or initialize them  
randomly.*

"""

**if** train\_from\_scratch==**False**:

    self.restore\_model()

*# Initialize best loss. This variable will store the lowest loss on the  
# eval dataset.*

best\_loss = 999

*# Initialize classes to update the mean loss of train and eval*

train\_loss = tf.keras.metrics.Mean('train\_loss')

eval\_loss = tf.keras.metrics.Mean('eval\_loss')

acc\_train = tf.keras.metrics.Mean('train\_acc')

acc\_val = tf.keras.metrics.Mean('val\_acc')

*# Initialize dictionary to store the loss history*

self.history = {}

self.history['train\_loss'] = []

self.history['eval\_loss'] = []

self.history['train\_acc'] = []

self.history['val\_acc'] = []

*# Begin training*

**with** tf.device(self.device):

**for** i **in** range(num\_epochs):

*# Training with gradient descent*

*#training\_data\_x = training\_data.shuffle(buffer\_size=1024).batch(1*

28)

**for** step, (images, target) **in** enumerate(training\_data):

            grads = self.grads\_fn(images, target, **True**)

            optimizer.apply\_gradients(zip(grads, self.variables))

*# Compute the loss on the training data after one epoch*

**for** step, (images, target) **in** enumerate(training\_data):

            loss = self.loss\_fn(images, target, **False**)

            accuracy = self.compute\_accuracy\_2(images, target)

            acc\_train(accuracy)

            train\_loss(loss)



```

self.history['train_loss'].append(train_loss.result().numpy())
self.history['train_acc'].append(acc_train.result().numpy())
# Reset metrics
train_loss.reset_states()
acc_train.reset_states()

# Compute the loss on the eval data after one epoch
for step, (images, target) in enumerate(eval_data):
    loss = self.loss_fn(images, target, False)
    accuracy = self.compute_accuracy_2(images, target)
    acc_val(accuracy)
    eval_loss(loss)
self.history['eval_loss'].append(eval_loss.result().numpy())
self.history['val_acc'].append(acc_val.result().numpy())
# Reset metrics
eval_loss.reset_states()
acc_val.reset_states()

# Print train and eval losses
if (i==0) | ((i+1)%verbose==0):
    print('Train loss at epoch %d: ' % (i+1), self.history['train_loss'][-1])
    print('Train Acc at epoch %d: ' % (i+1), self.history['train_acc'][-1])
    print('Eval loss at epoch %d: ' % (i+1), self.history['eval_loss'][-1])
    print('Eval Acc at epoch %d: ' % (i+1), self.history['val_acc'][-1])

# Check for early stopping
if self.history['eval_loss'][-1]<best_loss:
    best_loss = self.history['eval_loss'][-1]
    count = early_stopping_rounds
else:
    count -= 1
if count==0:
    break

```

In [5]:

```

# Specify the path where you want to save/restore the trained variables.
checkpoint_directory = 'models_checkpoints/fmnist/'

# Use the GPU if available.
device = 'gpu:0'

# Define optimizer.
optimizer = tf.compat.v1.train.AdamOptimizer(learning_rate=1e-4)

# Instantiate model. This doesn't initialize the variables yet.
model = ImageRecognitionCNN(num_classes=10, device=device,
                             checkpoint_directory=checkpoint_directory)

#model = ImageRecognitionCNN(num_classes=7, device=device)

```

In [6]:

```
# Train model
model.fit_fc(train_dataset, val_dataset, optimizer, num_epochs=10,
              early_stopping_rounds=2, verbose=2, train_from_scratch=True)
```

```
2022-03-22 19:56:38.947088: I tensorflow/stream_executor/cuda/cuda_dnn
n.cc:368] Loaded cuDNN version 8101
2022-03-22 19:56:39.094496: W tensorflow/stream_executor/gpu/asm_compil
er.cc:111] *** WARNING *** You are using ptxas 11.0.194, which is old
er than 11.1. ptxas before 11.1 is known to miscompile XLA code, leadi
ng to incorrect results or invalid-address errors.
```

You may not need to update to CUDA 11.1; cherry-picking the ptxas binary is often sufficient.

```
Train loss at epoch 1: 0.41828042
Train Acc at epoch 1: 0.8567134
Eval loss at epoch 1: 0.4425087
Eval Acc at epoch 1: 0.85722053
Train loss at epoch 2: 0.32595122
Train Acc at epoch 2: 0.864509
Eval loss at epoch 2: 0.35622242
Eval Acc at epoch 2: 0.87050176
Train loss at epoch 4: 0.25759926
Train Acc at epoch 4: 0.8836127
Eval loss at epoch 4: 0.3094024
Eval Acc at epoch 4: 0.88693607
Train loss at epoch 6: 0.21345654
Train Acc at epoch 6: 0.8951838
Eval loss at epoch 6: 0.28155872
Eval Acc at epoch 6: 0.89767534
Train loss at epoch 8: 0.18166418
Train Acc at epoch 8: 0.90346515
Eval loss at epoch 8: 0.26650384
Eval Acc at epoch 8: 0.9054872
Train loss at epoch 10: 0.15592033
Train Acc at epoch 10: 0.9106455
Eval loss at epoch 10: 0.2593014
Eval Acc at epoch 10: 0.91233075
```

In [ ]: