

In [ ]:

```
import os
import numpy as np
import time
import random
import tensorflow as tf
from tensorflow import keras

import matplotlib.pyplot as plt
```

```
np.random.seed(1234)
tf.random.set_seed(1234)
```

In [ ]:

```
tf.config.list_physical_devices('GPU')
```

In [ ]:

```
(X_train, y_train), (X_test, y_test) = keras.datasets.fashion_mnist.load_data() #
Load MNIST or FMNIST
assert X_train.shape == (60000, 28, 28)
assert X_test.shape == (10000, 28, 28)
assert y_train.shape == (60000,)
assert y_test.shape == (10000,)

# Display randomly selected data
indices = list(np.random.randint(X_train.shape[0], size=3))
for i in range(3):
    plt.subplot(1,3,i+1)
    plt.imshow(X_train[indices[i]].reshape(28,28), cmap='gray', interpolation='none')
    plt.title("Index {} Class {}".format(indices[i], y_train[indices[i]]))
plt.tight_layout()
```

In [ ]:

```
# Split train dataset into train and validation
X_val  = X_train[50000:60000]
X_train = X_train[0:50000]
y_val  = y_train[50000:60000]
y_train = y_train[0:50000]

print("size of training set is", str(X_train.shape[0]), "samples")
print("every train example is", str(X_train.shape[1]), "by", str(X_train.shape[2]
))

print("size of validation set is", str(X_val.shape[0]), "samples")
print("every validation example is", str(X_val.shape[1]), "by", str(X_val.shape[2]
))

X_train = X_train.reshape(50000, 28*28)
X_val = X_val.reshape(10000, 28*28)
X_test = X_test.reshape(10000, 28*28)

print("size of training set is", str(X_train.shape[0]), "samples")
print("every train example has", str(X_train.shape[1]), "features")

print("size of validation set is", str(X_val.shape[0]), "samples")
print("every validation example has", str(X_val.shape[1]), "features")

# Split dataset into batches
#train_ds = tf.data.Dataset.from_tensor_slices((X_train, y_train)).batch(16)
#test_ds = tf.data.Dataset.from_tensor_slices((X_test, y_test)).batch(4)
```

In [ ]:

```
#Normalize Data

X_train = X_train/255
X_val = X_val/255
X_test = X_test/255

print(f'max: ', np.max(X_train))
print(f'min: ', np.min(X_train))
```

In [ ]:

```
size_input = X_train.shape[1]

size_hidden1 = 128
size_hidden2 = 128
size_hidden3 = 128

size_output = 10

number_of_train_examples = X_train.shape[0]
number_of_test_examples = X_test.shape[0]

y_train = tf.keras.utils.to_categorical(y_train, num_classes=10) # Other function
is tf.one_hot(y_train, depth=10)
y_val = tf.keras.utils.to_categorical(y_val, num_classes=10)
y_test = tf.keras.utils.to_categorical(y_test, num_classes=10)

print(tf.shape(y_val))
```

In [ ]:

```

# Define class to build mlp model
class MLP(object):
    def __init__(self, size_input, size_hidden1, size_hidden2, size_hidden3, size_output, device=None):
        """
        size_input: int, size of input layer
        size_hidden1: int, size of the 1st hidden layer
        size_hidden2: int, size of the 2nd hidden layer
        size_output: int, size of output layer
        device: str or None, either 'cpu' or 'gpu' or None. If None, the device to
        be used will be decided automatically during Eager Execution
        """
        self.size_input, self.size_hidden1, self.size_hidden2, self.size_hidden3,
self.size_output, self.device = \
            size_input, size_hidden1, size_hidden2, size_hidden3, size_output, device

        # Initialize weights between input mapping and a layer  $g(f(x)) = \text{layer}$ 
        self.W1 = tf.Variable(tf.random.normal([self.size_input, self.size_hidden1], stddev=0.1)) # Xavier(Fan-in fan-out) and Orthogonal
        # Initialize biases for hidden layer
        self.b1 = tf.Variable(tf.zeros([1, self.size_hidden1])) # 0 or constant(0.01)

        # Initialize weights between input layer and 1st hidden layer
        self.W2 = tf.Variable(tf.random.normal([self.size_hidden1, self.size_hidden2], stddev=0.1))
        # Initialize biases for hidden layer
        self.b2 = tf.Variable(tf.zeros([1, self.size_hidden2]))

        # Initialize weights between 1st hidden layer and 2nd hidden layer
        self.W3 = tf.Variable(tf.random.normal([self.size_hidden2, self.size_hidden3], stddev=0.1))
        # Initialize biases for hidden layer
        self.b3 = tf.Variable(tf.zeros([1, self.size_hidden3]))

        # Initialize weights between 2nd hidden layer and output layer
        self.W4 = tf.Variable(tf.random.normal([self.size_hidden3, self.size_output], stddev=0.1))
        # Initialize biases for output layer
        self.b4 = tf.Variable(tf.zeros([1, self.size_output]))

        # Initializing gamma and beta for all the layers for batch norm
        self.gamma = {"1": tf.Variable(tf.ones(self.W1.shape[-1])),
                      "2": tf.Variable(tf.ones(self.W2.shape[-1])),
                      "3": tf.Variable(tf.ones(self.W3.shape[-1])),
                      "4": tf.Variable(tf.ones(self.W4.shape[-1]))}
        self.beta = {"1": tf.Variable(tf.zeros(self.W1.shape[-1])),
                     "2": tf.Variable(tf.zeros(self.W2.shape[-1])),
                     "3": tf.Variable(tf.zeros(self.W3.shape[-1])),
                     "4": tf.Variable(tf.zeros(self.W4.shape[-1]))}

        # Define variables to be updated during backpropagation

```

```

self.variables = [self.W1, self.W2, self.W3, self.W4,
                  self.b1, self.b2, self.b3, self.b4,
                  self.gamma["1"], self.beta["1"],
                  self.gamma["2"], self.beta["2"],
                  self.gamma["3"], self.beta["3"]]
                  #self.gamma["4"], self.beta["4"]]

self.mean = {"1": tf.Variable(tf.zeros(self.W1.shape[-1])),
             "2": tf.Variable(tf.zeros(self.W2.shape[-1])),
             "3": tf.Variable(tf.zeros(self.W3.shape[-1])),
             #"4": tf.Variable(tf.zeros(self.W4.shape[-1])),
             }
self.var = {"1": tf.Variable(tf.zeros(self.W1.shape[-1])),
            "2": tf.Variable(tf.zeros(self.W2.shape[-1])),
            "3": tf.Variable(tf.zeros(self.W3.shape[-1])),
            #"4": tf.Variable(tf.zeros(self.W4.shape[-1])),
            }

def forward(self, X, run):
    """
    forward pass
    X: Tensor, inputs
    """
    if self.device is not None:
        with tf.device('gpu:0' if self.device=='gpu' else 'cpu'):
            if run == "train":
                self.y = self.compute_output_train(X)
            else:
                self.y = self.compute_output_test(X)
    else:
        if run == "train":
            self.y = self.compute_output_train(X)
        else:
            self.y = self.compute_output_test(X)

    return self.y

def loss(self, y_pred, y_true):
    """
    y_pred - Tensor of shape (batch_size, size_output)
    y_true - Tensor of shape (batch_size, size_output)
    """
    #y_true_tf = tf.cast(tf.reshape(y_true, (-1, self.size_output)), dtype=tf.
float32)
    y_true_tf = tf.cast(y_true, dtype=tf.float32)
    y_pred_tf = tf.cast(y_pred, dtype=tf.float32)
    cce = tf.keras.losses.CategoricalCrossentropy(from_logits=True)
    loss_x = cce(y_true_tf, y_pred_tf)
    # Use keras or tf_softmax, both should work for any given model
    #loss_x = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y_
pred_tf, labels=y_true_tf))

    return loss_x

def backward(self, X_train, y_train, opti):
    """

```

```

    backward pass
    """
    optimizer = opti

    with tf.GradientTape() as tape:
        predicted = self.forward(X_train, "train")
        current_loss = self.loss(predicted, y_train)

    grads = tape.gradient(current_loss, self.variables)
    optimizer.apply_gradients(zip(grads, self.variables))

def batch_norm(self, inputs, layer):
    mean, var = tf.nn.moments(inputs, [0])
    self.mean[layer] = 0.9*self.mean[layer] + 0.1*mean
    self.var[layer] = 0.9*self.var[layer] + 0.1*var

    return tf.add(tf.divide(tf.multiply(self.gamma[layer], tf.subtract(inputs,
mean)), tf.sqrt(var+1e-8)), self.beta[layer])

def compute_output_train(self, X):
    """
    Custom method to obtain output tensor during forward pass
    """
    # Cast X to float32
    X_tf = tf.cast(X, dtype=tf.float32)
    #X_tf = X

    # Compute values in hidden layers
    z1 = tf.matmul(X_tf, self.W1) + self.b1
    z1 = self.batch_norm(z1, "1")
    h1 = tf.nn.relu(z1)

    z2 = tf.matmul(h1, self.W2) + self.b2
    z2 = self.batch_norm(z2, "2")
    h2 = tf.nn.relu(z2)

    z3 = tf.matmul(h2, self.W3) + self.b3
    z3 = self.batch_norm(z3, "3")
    h3 = tf.nn.relu(z3)

    # Compute output
    output = tf.matmul(h3, self.W4) + self.b4
    #output = self.batch_norm(output, "4")
    #Now consider two things , First look at inbuild loss functions if they wo
rk with softmax or not and then change this
    # Second add tf.Softmax(output) and then return this variable
    return (output)

def compute_output_test(self, X):
    """
    Custom method to obtain output tensor during forward pass
    """
    # Cast X to float32
    X_tf = tf.cast(X, dtype=tf.float32)
    #X_tf = X

```

```
# Compute values in hidden layers
z1 = tf.matmul(X_tf, self.W1) + self.b1
z1 = tf.add(tf.divide(tf.multiply(self.gamma["1"], tf.subtract(z1, self.mean["1"])), tf.sqrt(self.var["1"])), self.beta["1"])
h1 = tf.nn.relu(z1)

z2 = tf.matmul(h1, self.W2) + self.b2
z2 = tf.add(tf.divide(tf.multiply(self.gamma["2"], tf.subtract(z2, self.mean["2"])), tf.sqrt(self.var["2"])), self.beta["2"])
h2 = tf.nn.relu(z2)

z3 = tf.matmul(h2, self.W3) + self.b3
z3 = tf.add(tf.divide(tf.multiply(self.gamma["3"], tf.subtract(z3, self.mean["3"])), tf.sqrt(self.var["3"])), self.beta["3"])
h3 = tf.nn.relu(z3)

# Compute output
output = tf.matmul(h3, self.W4) + self.b4

#Now consider two things , First look at inbuild loss functions if they work with softmax or not and then change this
# Second add tf.Softmax(output) and then return this variable
return (output)
```

In [10]:

```

seeds = random.sample(range(1000, 9999), 3)
for trail in range(3):

    seed = seeds[trail]

    print(f'\n\n***** Running MLP with pre-activation Batch Normalization for seed {seed} *****\n')
    # Set number of epochs
    NUM_EPOCHS = 10

    # Initialize model using CPU
    mlp_on_gpu = MLP(size_input, size_hidden1, size_hidden2, size_hidden3, size_output, device='gpu')

    time_start = time.time()
    opti = tf.keras.optimizers.SGD(learning_rate = 0.1)

    for epoch in range(NUM_EPOCHS):
        loss_total = tf.zeros([1,1], dtype=tf.float32)
        lt = 0

        train_ds = tf.data.Dataset.from_tensor_slices((X_train, y_train)).shuffle(
25, seed=epoch*(seed)).batch(128)
        kz = 0
        accuracy_z = 0.0
        cur_train_acc = 0.0
        for inputs, outputs in train_ds:
            qw, tr = tf.shape(inputs)
            kz = kz + 1
            preds = mlp_on_gpu.forward(inputs, "train")
            loss_total = loss_total + mlp_on_gpu.loss(preds, outputs)
            lt = lt + mlp_on_gpu.loss(preds, outputs)
            mlp_on_gpu.backward(inputs, outputs, opti)

            preds = mlp_on_gpu.forward(X_train, "train")
            # Get probs, remember we only have logits from our forward function, we need to apply softmax on top of it to get probs
            preds = tf.nn.softmax(preds)
            correct_prediction = tf.equal(tf.argmax(preds, 1), tf.argmax(y_train, 1))
            accuracy_z = accuracy_z + tf.reduce_mean(tf.cast(correct_prediction, "float"))

        cur_train_acc += accuracy_z.numpy()
        ds = cur_train_acc
        print('\nTrain Accuracy: {:.4f}'.format(ds))
        print('Number of Epoch = {} - Average Cross Entropy:= {} '.format(epoch +
1, np.sum(loss_total) / X_train.shape[0]))
        preds_val = mlp_on_gpu.forward(X_val, "train")
        preds_val = tf.nn.softmax(preds_val)
        correct_prediction = tf.equal(tf.argmax(preds_val, 1), tf.argmax(y_val, 1))

    # Calculate accuracy
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))

```



```
cur_val_acc = accuracy.numpy()

print('\nValidation Accuracy: {:.4f}'.format(cur_val_acc))

plt.plot(epoch + 1, np.sum(loss_total) / X_train.shape[0], 'go')

time_taken = time.time() - time_start
plt.show()
# Validate model

print('\nTotal time taken (in seconds): {:.2f}'.format(time_taken))
#For per epoch_time = Total_Time / Number_of_epochs

test_loss_total = tf.Variable(0, dtype=tf.float32)
correct_prediction = tf.Variable(0, dtype=tf.float32)

test_ds = tf.data.Dataset.from_tensor_slices((X_test, y_test)).batch(4)

#test_loss_total = 0.0
for inputs, outputs in test_ds:
    preds = mlp_on_gpu.forward(inputs, "test")
    test_loss_total = test_loss_total + mlp_on_gpu.loss(preds, outputs)
print('Test loss: {:.4f}'.format(np.sum(test_loss_total.numpy()) / X_test.shape[0]))

# Test model
preds_test = mlp_on_gpu.forward(X_test, "test")
preds_test = tf.nn.softmax(preds_test)
correct_prediction = tf.equal(tf.argmax(preds_test, 1), tf.argmax(y_test, 1))

# Calculate accuracy
accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
cur_test_acc = accuracy.numpy()
print('\nTest Accuracy: {:.2f}'.format(cur_test_acc))
```

\*\*\*\*\* Running MLP with pre-activation Batch Normalization  
for seed 2977 \*\*\*\*\*

Train Accuracy: 0.8617  
Number of Epoch = 1 - Average Cross Entropy:= 0.004356533813476562

Validation Accuracy: 0.8495

Train Accuracy: 0.8788  
Number of Epoch = 2 - Average Cross Entropy:= 0.0029947314453125

Validation Accuracy: 0.8659

Train Accuracy: 0.8902  
Number of Epoch = 3 - Average Cross Entropy:= 0.002642081604003906

Validation Accuracy: 0.8721

Train Accuracy: 0.9009  
Number of Epoch = 4 - Average Cross Entropy:= 0.0024121412658691405

Validation Accuracy: 0.8759

Train Accuracy: 0.9096  
Number of Epoch = 5 - Average Cross Entropy:= 0.002219288635253906

Validation Accuracy: 0.8788

Train Accuracy: 0.9147  
Number of Epoch = 6 - Average Cross Entropy:= 0.002058863525390625

Validation Accuracy: 0.8795

Train Accuracy: 0.9211  
Number of Epoch = 7 - Average Cross Entropy:= 0.001903658447265625

Validation Accuracy: 0.8822

Train Accuracy: 0.9256  
Number of Epoch = 8 - Average Cross Entropy:= 0.0017605203247070312


Validation Accuracy: 0.8800

Train Accuracy: 0.9294  
Number of Epoch = 9 - Average Cross Entropy:= 0.00164457275390625

Validation Accuracy: 0.8785

Train Accuracy: 0.9341  
Number of Epoch = 10 - Average Cross Entropy:= 0.0015227243041992188

Validation Accuracy: 0.8782



Total time taken (in seconds): 107.32

Test loss: 0.0958

Test Accuracy: 0.87

\*\*\*\*\* Running MLP with pre-activation Batch Normalization  
for seed 1427 \*\*\*\*\*

Train Accuracy: 0.8587

Number of Epoch = 1 - Average Cross Entropy:= 0.00438836181640625

Validation Accuracy: 0.8476

Train Accuracy: 0.8780

Number of Epoch = 2 - Average Cross Entropy:= 0.0030348040771484377

Validation Accuracy: 0.8601

Train Accuracy: 0.8888

Number of Epoch = 3 - Average Cross Entropy:= 0.0026650567626953126

Validation Accuracy: 0.8650

Train Accuracy: 0.9002

Number of Epoch = 4 - Average Cross Entropy:= 0.0024279241943359377

Validation Accuracy: 0.8714

Train Accuracy: 0.9047

Number of Epoch = 5 - Average Cross Entropy:= 0.0022358851623535156

Validation Accuracy: 0.8710

Train Accuracy: 0.9112

Number of Epoch = 6 - Average Cross Entropy:= 0.002071214599609375

Validation Accuracy: 0.8748

Train Accuracy: 0.9190

Number of Epoch = 7 - Average Cross Entropy:= 0.0019120587158203124

Validation Accuracy: 0.8759

Train Accuracy: 0.9236

Number of Epoch = 8 - Average Cross Entropy:= 0.0017844329833984374

Validation Accuracy: 0.8789

Train Accuracy: 0.9290


Number of Epoch = 9 - Average Cross Entropy:= 0.0016624235534667968

Validation Accuracy: 0.8817

Train Accuracy: 0.9308

Number of Epoch = 10 - Average Cross Entropy:= 0.0015360650634765624

Validation Accuracy: 0.8767



Total time taken (in seconds): 115.79

Test loss: 0.1016

Test Accuracy: 0.87

\*\*\*\*\* Running MLP with pre-activation Batch Normalization  
for seed 9846 \*\*\*\*\*

Train Accuracy: 0.8614

Number of Epoch = 1 - Average Cross Entropy:= 0.004333209228515625

Validation Accuracy: 0.8449

Train Accuracy: 0.8805

Number of Epoch = 2 - Average Cross Entropy:= 0.0029972756958007813

Validation Accuracy: 0.8608

Train Accuracy: 0.8918

Number of Epoch = 3 - Average Cross Entropy:= 0.0026356414794921875

Validation Accuracy: 0.8692

Train Accuracy: 0.9004

Number of Epoch = 4 - Average Cross Entropy:= 0.0024016685485839843

Validation Accuracy: 0.8744

Train Accuracy: 0.9084

Number of Epoch = 5 - Average Cross Entropy:= 0.002214339141845703

Validation Accuracy: 0.8771

Train Accuracy: 0.9153

Number of Epoch = 6 - Average Cross Entropy:= 0.0020412139892578124

Validation Accuracy: 0.8782

Train Accuracy: 0.9199

Number of Epoch = 7 - Average Cross Entropy:= 0.0019056634521484375

Validation Accuracy: 0.8800

Train Accuracy: 0.9246

Number of Epoch = 8 - Average Cross Entropy:= 0.0017654421997070312

Validation Accuracy: 0.8798

Train Accuracy: 0.9310


Number of Epoch = 9 - Average Cross Entropy:= 0.0016393472290039063

Validation Accuracy: 0.8785

Train Accuracy: 0.9326

Number of Epoch = 10 - Average Cross Entropy:= 0.0015347804260253907

Validation Accuracy: 0.8825



Total time taken (in seconds): 111.65

Test loss: 0.0923

Test Accuracy: 0.88

In [ ]: