

In [35]:

```
import os
import copy
import time
import shelve
import random
import numpy as np
import pandas as pd
import tensorflow as tf
from pandas import DataFrame
import matplotlib.pyplot as plt
```

In [36]:

```
def load_data(name):  
    if name == 'mnist':  
        (X_train, y_train), (X_test, y_test) = tf.keras.datasets.mnist.load_data()  
    elif name == 'fashion_mnist':  
        (X_train, y_train), (X_test, y_test) = tf.keras.datasets.fashion_mnist.load_data()  
    else:  
        print('Only mnist or fashion_mnist.')  
        return False  
  
    imageSize = X_train.shape[1]*X_train.shape[2]  
    numClasses = np.max(y_train)+1  
  
    X_train = np.reshape(X_train.astype(float)/255.0, (-1, 784))  
    X_test = np.reshape(X_test.astype(float)/255.0, (-1, 784))  
  
    y_train = tf.keras.utils.to_categorical(y_train, num_classes=numClasses)  
    y_test = tf.keras.utils.to_categorical(y_test, num_classes=numClasses)  
  
    X_val = X_train[-10000:]  
    y_val = y_train[-10000:]  
    X_train = X_train[:-10000]  
    y_train = y_train[:-10000]  
  
    print('Data Split: ')  
    print(f'X_train: {X_train.shape}, y_train: {y_train.shape}')  
    print(f'X_test : {X_test.shape }, y_test : {y_test.shape }')  
    print(f'X_val  : {X_val.shape }, y_val  : {y_val.shape }')  
  
    data = {}  
    data['X_train'] = X_train  
    data['y_train'] = y_train  
    data['X_val'] = X_val  
    data['y_val'] = y_val  
    data['X_test'] = X_test  
    data['y_test'] = y_test  
  
    data['imageSize'] = imageSize  
  
    return data
```

In [37]:

```

class myOptimizer():

    def __init__(self, var):

        self.m = [np.zeros(np.shape(i)) for i in var]
        self.v = [np.zeros(np.shape(i)) for i in var]
        self.u = [np.zeros(np.shape(i)) for i in var]

        self.alpha    = 5e-4
        self.beta      = [0.9, 0.999, 0.999987]
        self.epsilon   = [1e-8, 1e-6]

    def updateWeights(self, t, grads, var):

        for i in range(np.shape(var)[0]):

            self.m[i] = (self.beta[0]*self.m[i]) + ((1 - self.beta[0])*grads[i])
            self.v[i] = (self.beta[1]*self.v[i]) + ((1 - self.beta[1])*(grads[i]*g
rads[i]))
            self.u[i] = (self.beta[2]*self.u[i]) + ((1 - self.beta[2])*(grads[i]*g
rads[i]*grads[i]))

            mHat = self.m[i]/(1-self.beta[0]**t)
            vHat = self.v[i]/(1-self.beta[1]**t)
            uHat = self.u[i]/(1-self.beta[2]**t)

            # delta = self.alpha*mHat/((tf.sign(vHat)*tf.sqrt(tf.abs(vHat)))+(tf.s
            ign(uHat)*tf.math.pow(tf.abs(uHat), 1/3)*self.epsilon[0])+self.epsilon[0])
            delta = self.alpha * mHat / (tf.sqrt(vHat)+(tf.sign(uHat)*tf.math.pow(
            tf.abs(uHat), 1/3)*self.epsilon[0] + self.epsilon[1]))
            var[i].assign(var[i] - delta)

```

In [38]:

```

class MLP(object):

    def __init__(self, name, size_input, size_hidden, size_output, learning_rate=
0.01, weight_coeff=1,\
                    Reg=None, RegC=0, training=None, validation=None, accuracy=0, de
vice=None):

        self.name            = name
        self.size_input      = size_input
        self.size_hidden     = size_hidden
        self.size_output     = size_output
        self.Reg             = Reg
        self.RegC            = RegC
        self.training        = training
        self.validation       = validation
        self.accuracy        = accuracy
        self.device          = device
        self.learning_rate   = learning_rate
        self.weight_coeff    = weight_coeff

        self.W1 = self.initWeights(self.size_input, self.size_hidden[0], self.weig
ht_coeff)
        self.b1 = self.initWeights(1, self.size_hidden[0], self.weight_coeff)

        self.W2 = self.initWeights(self.size_hidden[0], self.size_hidden[1], self.
weight_coeff)
        self.b2 = self.initWeights(1, self.size_hidden[1], self.weight_coeff)

        self.W3 = self.initWeights(self.size_hidden[1], self.size_hidden[2], self.
weight_coeff)
        self.b3 = self.initWeights(1, self.size_hidden[2], self.weight_coeff)

        self.W4 = self.initWeights(self.size_hidden[2], self.size_output, self.wei
ght_coeff)
        self.b4 = self.initWeights(1, self.size_output, self.weight_coeff)

        self.variables = [self.W1, self.b1, self.W2, self.b2, self.W3, self.b3, se
lf.W4, self.b4]

        self.optimizer = myOptimizer(self.variables)

    def initWeights(self, rows, columns, multFactor=1):
        return tf.Variable(multFactor*tf.random.normal([rows, columns]))

    def forward(self, X):

        if self.device is not None:
            with tf.device('gpu:0' if self.device=='gpu' else 'cpu'):
                self.y = self.compute_output(X)
        else:
            self.y = self.compute_output(X)

        return self.y

```

```

def getRegLoss(self, X_train):

    if self.Reg=='L2':
        return (self.RegC/X_train.shape[0])*(tf.reduce_sum(tf.math.square(self
        .variables[0])) +
                                                    tf.reduce_sum(tf.math.square(self
        .variables[2])) +
                                                    tf.reduce_sum(tf.math.square(self
        .variables[4])) +
                                                    tf.reduce_sum(tf.math.square(self
        .variables[6])))

    elif self.Reg=='L1':
        return (self.RegC/X_train.shape[0])*tf.abs(tf.reduce_sum(self.variable
s[0]) +
                                                    tf.reduce_sum(self.variable
s[2]) +
                                                    tf.reduce_sum(self.variable
s[4]) +
                                                    tf.reduce_sum(self.variable
s[6]))

    elif self.Reg=='L1+L2':
        L2 = (self.RegC/X_train.shape[0])*(tf.reduce_sum(tf.math.square(self.
variables[0])) +
                                                    tf.reduce_sum(tf.math.square(self.
variables[2])) +
                                                    tf.reduce_sum(tf.math.square(self.
variables[4])) +
                                                    tf.reduce_sum(tf.math.square(self.
variables[6])))

        L1 = (self.RegC/X_train.shape[0])*tf.abs(tf.reduce_sum(self.variables[
0]) +
                                                    tf.reduce_sum(self.variables[
2]) +
                                                    tf.reduce_sum(self.variables[
4]) +
                                                    tf.reduce_sum(self.variables[
6]))

        return L1+L2

    else:
        return 0

def loss(self, y_pred, y_true):

    y_true_tf = tf.cast(tf.reshape(y_true, (-1, self.size_output)), dtype=tf.f
loat32)
    y_pred_tf = tf.cast(y_pred, dtype=tf.float32)

    loss = tf.keras.losses.CategoricalCrossentropy()(y_true_tf, y_pred_tf)
    return loss

def backward(self, X_train, y_train, t):

```

```
with tf.GradientTape() as tape:
    predicted = self.forward(X_train)
    current_loss = self.loss(predicted, y_train)
    current_loss += self.getRegLoss(X_train)

grads = tape.gradient(current_loss, self.variables)
self.optimizer.updateWeights(t, grads, self.variables)

def compute_output(self, X):
    X_tf = tf.cast(X, dtype=tf.float32)

    w1Hat = tf.matmul(X_tf, self.variables[0]) + self.variables[1]
    h1Hat = tf.nn.relu(w1Hat)

    w2Hat = tf.matmul(h1Hat, self.variables[2]) + self.variables[3]
    h2Hat = tf.nn.relu(w2Hat)

    w3Hat = tf.matmul(h2Hat, self.variables[4]) + self.variables[5]
    h3Hat = tf.nn.relu(w3Hat)

    w4Hat = tf.matmul(h3Hat, self.variables[6]) + self.variables[7]
    output = tf.nn.softmax(w4Hat)

    return output

def getAccuracy(self, predictions, outputs):
    preds = np.argmax(predictions, axis=1)
    y_true = np.argmax(outputs, axis=1)

    return (preds==y_true).mean()
```

In [39]:

```

def trainModel(model, data, NUM_EPOCHS=10, batchSize=50, seedVal=1234):

    X_train = data['X_train']
    y_train = data['y_train']
    X_val    = data['X_val']
    y_val    = data['y_val']

    training = np.zeros(shape=(NUM_EPOCHS, 3))
    validation = np.zeros(shape=(NUM_EPOCHS, 3))

    train_ds = tf.data.Dataset.from_tensor_slices((X_train, y_train)).batch(batchSize)
    val_ds    = tf.data.Dataset.from_tensor_slices((X_val, y_val)).batch(batchSize)

    print(f'\n\n***** Training model: {model.name} with seed: {seedVal} *****\n')
    time_start = time.time()
    for epoch in range(NUM_EPOCHS):
        train_loss = tf.zeros([1, 1], dtype=tf.float32)
        val_loss    = tf.zeros([1, 1], dtype=tf.float32)

        train_ds = tf.data.Dataset.from_tensor_slices((X_train, y_train)).shuffle(
25, seed = epoch*(seedVal)).batch(batchSize)
        val_ds    = tf.data.Dataset.from_tensor_slices((X_val, y_val)).shuffle(25,
seed = epoch*(seedVal)).batch(batchSize)

        t = 0
        for inputs, outputs in train_ds:
            t+=1
            train_pred = model.forward(inputs)
            train_loss = train_loss + model.loss(train_pred, outputs)
            model.backward(inputs, outputs, t)
            train_acc = model.getAccuracy(train_pred, outputs)

        for inputs, outputs in val_ds:
            val_pred = model.forward(inputs)
            val_loss = val_loss + model.loss(val_pred, outputs)
            val_acc = model.getAccuracy(val_pred, outputs)

        # train_loss = np.array(train_loss)
        # val_loss = np.array(val_loss)

        training[epoch] = [epoch+1, train_acc, np.sum(train_loss)/X_train.shape[0]
]]
        validation[epoch] = [epoch+1, val_acc, np.sum(val_loss)/X_val.shape[0]
]]

    print('# Epoch:={}/{} - train loss:={:.4f} - val loss:={:.4f}, train acc:
={:.2f} - val acc:={:.2f}'\
          .format(epoch+1, NUM_EPOCHS, np.sum(train_loss)/X_train.shape[0], np
.sum(val_loss)/X_val.shape[0], train_acc, val_acc))

    time_taken = time.time()-time_start

```

```
print(f'\nTotal time taken (in seconds): {time_taken: .2f}')
print(f'\nFinished training model: {model.name}\n')
model.training = training
model.validation = validation

def testModel(model, data):

    X_test = data['X_test']
    y_test = data['y_test']

    preds = model.forward(X_test)

    pred = np.argmax(preds, axis=1)
    y_true= np.argmax(y_test, axis=1)

    model.accuracy = (pred==y_true).mean()*100

    print(f'***** Testing *****')
    print(f'{model.name} model accuracy = {model.accuracy:.2f}%')
    print(f'*****')

def plotAccuracyAndLoss(model):

    training = model.training
    validation = model.validation
    fig, (ax1, ax2) = plt.subplots(1, 2)
    training[:, -1] = training[:, -1]/np.linalg.norm(training[:, -1])
    ax1.plot(training[:,0], training[:,1], 'g')
    ax1.plot(training[:,0], training[:,2], 'b')
    ax1.set_title('Training')
    ax1.legend(["Accuracy", "Loss"])

    validation[:, -1] = validation[:, -1]/np.linalg.norm(validation[:, -1])
    ax2.plot(validation[:,0], validation[:,1], 'g')
    ax2.plot(validation[:,0], validation[:,2], 'b')
    ax2.set_title('Validation')
    ax2.legend(["Accuracy", "Loss"])
    plt.show()
```


In [44]:

```
def main():  
    data = load_data('mnist')  
    imageSize = data['imageSize']  
  
    size_input = imageSize  
    size_hidden = [512, 256, 128]  
    size_output = 10  
  
    numEpochs = 15  
    batchSize = 50  
  
    learning_rate = 1e-4  
    weight_coeff = 0.1  
  
    seeds = np.random.randint(1000, 9999)  
  
    np.random.seed(seeds)  
    tf.random.set_seed(seeds)  
  
    model = MLP('mlp_on_gpu_default', size_input, size_hidden, size_output, learning_rate=learning_rate,\  
                weight_coeff=weight_coeff, device='gpu')  
  
    trainModel(model, data, numEpochs, batchSize, seeds)  
    testModel(model, data)  
    plotAccuracyAndLoss(model)  
  
    return model
```

In [45]:

```
if __name__ == "__main__":  
    mnist = main()
```

Data Split:

X_train: (50000, 784), y_train: (50000, 10)

X_test : (10000, 784), y_test : (10000, 10)

X_val : (10000, 784), y_val : (10000, 10)


***** Training model: mlp_on_gpu_default with seed: 4775 **

```
# Epoch:=1/15 - train loss:=0.0051 - val loss:=0.0025, train acc:=0.9
8 - val acc:=1.00
# Epoch:=2/15 - train loss:=0.0018 - val loss:=0.0019, train acc:=1.0
0 - val acc:=1.00
# Epoch:=3/15 - train loss:=0.0011 - val loss:=0.0018, train acc:=1.0
0 - val acc:=0.98
# Epoch:=4/15 - train loss:=0.0007 - val loss:=0.0018, train acc:=1.0
0 - val acc:=1.00
# Epoch:=5/15 - train loss:=0.0004 - val loss:=0.0018, train acc:=1.0
0 - val acc:=0.98
# Epoch:=6/15 - train loss:=0.0003 - val loss:=0.0021, train acc:=0.9
8 - val acc:=0.96
# Epoch:=7/15 - train loss:=0.0003 - val loss:=0.0020, train acc:=1.0
0 - val acc:=1.00
# Epoch:=8/15 - train loss:=0.0003 - val loss:=0.0021, train acc:=1.0
0 - val acc:=0.98
# Epoch:=9/15 - train loss:=0.0002 - val loss:=0.0021, train acc:=0.9
8 - val acc:=0.98
# Epoch:=10/15 - train loss:=0.0002 - val loss:=0.0022, train acc:=0.
98 - val acc:=1.00
# Epoch:=11/15 - train loss:=0.0002 - val loss:=0.0019, train acc:=1.
00 - val acc:=1.00
# Epoch:=12/15 - train loss:=0.0001 - val loss:=0.0027, train acc:=1.
00 - val acc:=1.00
# Epoch:=13/15 - train loss:=0.0002 - val loss:=0.0021, train acc:=0.
98 - val acc:=1.00
# Epoch:=14/15 - train loss:=0.0001 - val loss:=0.0024, train acc:=1.
00 - val acc:=1.00
# Epoch:=15/15 - train loss:=0.0001 - val loss:=0.0024, train acc:=1.
00 - val acc:=0.96
```

Total time taken (in seconds): 305.84

Finished training model: mlp_on_gpu_default

***** Testing *****
mlp_on_gpu_default model accuracy = 98.05%



In []: