

In [1]:


```
import tensorflow as tf
import pandas as pd
import pickle
import matplotlib.pyplot as plt
%matplotlib inline
tf.random.set_seed(1234)
```

In [2]:

```
import sys
sys.path.insert(1, '/data/')
from data_utils import parse_imdb_sequence
```

In [3]:

```
length_reviews = pickle.load(open('./data/length_reviews.pkl', 'rb'))
pd.DataFrame(length_reviews, columns=['Length reviews']).hist(bins=100, color='blue');
plt.grid(False);
```



In [4]:

```
full_dataset = tf.data.TFRecordDataset('./data/train.tfrecords')
full_dataset = full_dataset.shuffle(buffer_size=10000)

DATASET_SIZE = sum(1 for _ in full_dataset)
print(DATASET_SIZE)

train_size = int(0.85 * DATASET_SIZE)
val_size = int(0.15 * DATASET_SIZE)

train_dataset = full_dataset.take(train_size)
val_dataset = full_dataset.take(val_size)

train_dataset_size = sum(1 for _ in train_dataset)
print(train_dataset_size)

val_dataset_size = sum(1 for _ in val_dataset)
print(val_dataset_size)

train_dataset = train_dataset.map(parse_imdb_sequence).shuffle(buffer_size=10000)
train_dataset = train_dataset.padded_batch(512, padded_shapes=(None, [], []))

val_dataset = val_dataset.map(parse_imdb_sequence).shuffle(buffer_size=10000)
val_dataset = val_dataset.padded_batch(128, padded_shapes=(None, [], []))

test_dataset = tf.data.TFRecordDataset('./data/test.tfrecords')
test_dataset = test_dataset.map(parse_imdb_sequence).shuffle(buffer_size=10000)
test_dataset = test_dataset.padded_batch(512, padded_shapes=(None, [], []))
```

```
2022-03-28 23:07:48.281394: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-03-28 23:07:48.290323: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-03-28 23:07:48.290656: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-03-28 23:07:48.292315: I tensorflow/core/platform/cpu_feature_guard.cc:151] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
2022-03-28 23:07:48.293619: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-03-28 23:07:48.293960: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-03-28 23:07:48.294208: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-03-28 23:07:48.617958: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-03-28 23:07:48.618318: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-03-28 23:07:48.618579: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:936] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2022-03-28 23:07:48.618829: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1525] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 6980 MB memory: -> device: 0, name: NVIDIA GeForce GTX 1070, pci bus id: 0000:01:00.0, compute capability: 6.1
```

```
25000
21250
3750
```

In [5]:

```
# Read the word vocabulary  
word2idx = pickle.load(open('./data/word2idx.pkl', 'rb'))
```

In [6]:

```

class RNNModel(tf.keras.Model):
    def __init__(self, embedding_size=100, cell_size=64, dense_size=128,
                  num_classes=2, vocabulary_size=None, rnn_cell='lstm',
                  device='cpu:0', checkpoint_directory=None):
        ''' Define the parameterized layers used during forward-pass, the device
            where you would like to run the computation on and the checkpoint
            directory. Additionally, you can also modify the default size of the
            network.

            Args:
                embedding_size: the size of the word embedding.
                cell_size: RNN cell size.
                dense_size: the size of the dense layer.
                num_classes: the number of labels in the network.
                vocabulary_size: the size of the word vocabulary.
                rnn_cell: string, either 'lstm' or 'ugrnn'.
                device: string, 'cpu:n' or 'gpu:n' (n can vary). Default, 'cpu:0'.
                checkpoint_directory: the directory where you would like to save o

        ...
        restore a model.
        super(RNNModel, self).__init__()

        # Weights initializer function
        w_initializer = tf.compat.v1.keras.initializers.glorot_uniform()

        # Biases initializer function
        b_initializer = tf.zeros_initializer()

        # Initialize weights for word embeddings
        self.embeddings = tf.keras.layers.Embedding(vocabulary_size, embedding_size,
                                                    embeddings_initializer=w_initializer)

        # Dense layer initialization
        self.dense_layer = tf.keras.layers.Dense(dense_size, activation=tf.nn.relu,
                                                  kernel_initializer=w_initializer,
                                                  bias_initializer=b_initializer)

        # Predictions layer initialization
        self.pred_layer = tf.keras.layers.Dense(num_classes, activation=None,
                                                  kernel_initializer=w_initializer,
                                                  bias_initializer=b_initializer)

        # Basic LSTM cell
        if rnn_cell=='lstm':
            self.rnn_cell = tf.compat.v1.nn.rnn_cell.BasicLSTMCell(cell_size)
        # Else RNN cell
        else:
            self.rnn_cell = tf.compat.v1.nn.rnn_cell.BasicRNNCell(cell_size)

```

```

# Define the device
self.device = device

# Define the checkpoint directory
self.checkpoint_directory = checkpoint_directory

def predict(self, X, seq_length, is_training):
    """
    Predicts the probability of each class, based on the input sample.

    Args:
        X: 2D tensor of shape (batch_size, time_steps).
        seq_length: the length of each sequence in the batch.
        is_training: Boolean. Either the network is predicting in
                     training mode or not.
    """

    # Get the number of samples within a batch
    num_samples = tf.shape(X)[0]

    # Initialize LSTM cell state with zeros
    state = self.rnn_cell.zero_state(num_samples, dtype=tf.float32)

    # Get the embedding of each word in the sequence
    embedded_words = self.embeddings(X)

    # Unstack the embeddings
    unstacked_embeddings = tf.unstack(embedded_words, axis=1)

    # Iterate through each timestep and append the predictions
    outputs = []
    for input_step in unstacked_embeddings:
        output, state = self.rnn_cell(input_step, state)
        outputs.append(output)

    # Stack outputs to (batch_size, time_steps, cell_size)
    outputs = tf.stack(outputs, axis=1)

    # Extract the output of the last time step, of each sample
    idxs_last_output = tf.stack([tf.range(num_samples),
                                tf.cast(seq_length-1, tf.int32)], axis=1)
    final_output = tf.gather_nd(outputs, idxs_last_output)

    # Add dropout for regularization
    #dropped_output = tf.compat.v1.layers.Dropout(final_output, rate=0.3, training=is_training)

    # Pass the last cell state through a dense layer (ReLU activation)
    dense = self.dense_layer(final_output)

    # Compute the unnormalized log probabilities
    logits = self.pred_layer(dense)
    return logits

def loss_fn(self, X, y, seq_length, is_training):
    """ Defines the loss function used during

```

```

        training.
        """
        preds = self.predict(X, seq_length, is_training)
        loss = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=y, logits=preds)
        return loss

    def grads_fn(self, X, y, seq_length, is_training):
        """ Dynamically computes the gradients of the loss value
            with respect to the parameters of the model, in each
            forward pass.
        """
        with tf.GradientTape() as tape:
            loss = self.loss_fn(X, y, seq_length, is_training)
        return tape.gradient(loss, self.variables)

    def restore_model(self):
        """ Function to restore trained model.
        """
        with tf.device(self.device):
            # Run the model once to initialize variables
            dummy_input = tf.constant(tf.zeros((1,1)))
            dummy_length = tf.constant(1, shape=(1,))
            dummy_pred = self.predict(dummy_input, dummy_length, False)
            # Restore the variables of the model
            saver = tf.compat.v1.train.Saver(self.variables)
            saver.restore(tf.train.latest_checkpoint(
                self.checkpoint_directory))

    def save_model(self, global_step=0):
        """ Function to save trained model.
        """
        tf.compat.v1.train.Saver(self.variables).save(save_path=self.checkpoint_directory,
                                                    global_step=global_step)

    def fit(self, training_data, eval_data, test_data, optimizer, num_epochs=500,
            early_stopping_rounds=10, verbose=10, train_from_scratch=False):
        """ Function to train the model, using the selected optimizer and
            for the desired number of epochs. You can either train from scratch
            or load the latest model trained. Early stopping is used in order to
            mitigate the risk of overfitting the network.

        Args:
            training_data: the data you would like to train the model on.
                           Must be in the tf.data.Dataset format.
            eval_data: the data you would like to evaluate the model on.
                       Must be in the tf.data.Dataset format.
            optimizer: the optimizer used during training.
            num_epochs: the maximum number of iterations you would like to
                           train the model.
            early_stopping_rounds: stop training if the accuracy on the eval
                                   dataset does not increase after n epochs.
            verbose: int. Specify how often to print the loss value of the network.
            train_from_scratch: boolean. Whether to initialize variables of the
        """

```

e

*the last trained model or initialize them randomly.*

```

"""

if train_from_scratch==False:
    self.restore_model()

# Initialize best_acc. This variable will store the highest accuracy
# on the eval dataset.
best_acc = 0

# Initialize classes to update the mean accuracy of train and eval
train_acc = tf.keras.metrics.Accuracy('train_acc')
eval_acc = tf.keras.metrics.Accuracy('eval_acc')
test_acc = tf.keras.metrics.Accuracy('test_acc')
# Initialize dictionary to store the accuracy history
self.history = {}
self.history['train_acc'] = []
self.history['eval_acc'] = []
self.history['test_acc'] = []

# Begin training
with tf.device(self.device):
    for i in range(num_epochs):
        # Training with gradient descent
        for step, (X, y, seq_length) in enumerate(training_data):
            grads = self.grads_fn(X, y, seq_length, True)
            optimizer.apply_gradients(zip(grads, self.variables))

        # Check accuracy train dataset
        for step, (X, y, seq_length) in enumerate(training_data):
            logits = self.predict(X, seq_length, False)
            preds = tf.argmax(logits, axis=1)
            train_acc(preds, y)
        self.history['train_acc'].append(train_acc.result().numpy())
        # Reset metrics
        train_acc.reset_states()

        # Check accuracy eval dataset
        for step, (X, y, seq_length) in enumerate(eval_data):
            logits = self.predict(X, seq_length, False)
            preds = tf.argmax(logits, axis=1)
            eval_acc(preds, y)
        self.history['eval_acc'].append(eval_acc.result().numpy())
        # Reset metrics
        eval_acc.reset_states()

        # Print train and eval accuracy
        if (i==0) | ((i+1)%verbose==0):
            print('Train accuracy at epoch %d: ' % (i+1), self.history['train_acc'][-1])
            print('Eval accuracy at epoch %d: ' % (i+1), self.history['eval_acc'][-1])

        # Check for early stopping

```



```
if self.history['eval_acc'][-1]>best_acc:
    best_acc = self.history['eval_acc'][-1]
    count = early_stopping_rounds
else:
    count -= 1
if count==0:
    break

for step, (X, y, seq_length) in enumerate(test_data):
    logits = self.predict(X, seq_length, False)
    preds = tf.argmax(logits, axis=1)
    test_acc(preds, y)
self.history['test_acc'].append(test_acc.result().numpy())
print('Test accuracy: ', self.history['test_acc'][-1])
```

In [7]:

```
# Specify the path where you want to save/restore the trained variables.
checkpoint_directory = './models_checkpoints/ImdbRNN1/'

# Use the GPU if available.
device = 'gpu:0'

# Define optimizer.
optimizer = tf.compat.v1.train.AdamOptimizer(learning_rate=1e-4)

# Instantiate model. This doesn't initialize the variables yet.
lstm_model1 = RNNModel(vocabulary_size=len(word2idx), device=device,
                       checkpoint_directory=checkpoint_directory)

# Train model
lstm_model1.fit(train_dataset, val_dataset, test_dataset, optimizer, num_epochs=20,
               early_stopping_rounds=5, verbose=1, train_from_scratch=True)

#lstm_model.save_model()
checkpoint = tf.train.Checkpoint(lstm_model1)
save_path = checkpoint.save(checkpoint_directory)

#####

# Define optimizer.
optimizer = tf.compat.v1.train.AdamOptimizer(learning_rate=1e-4)

# Instantiate model. This doesn't initialize the variables yet.
ugrnn_model1 = RNNModel(vocabulary_size=len(word2idx), rnn_cell='ugrnn',
                        device=device, checkpoint_directory=checkpoint_directory)

# Train model
ugrnn_model1.fit(train_dataset, val_dataset, test_dataset, optimizer, num_epochs=
20,
               early_stopping_rounds=5, verbose=1, train_from_scratch=True)

#lstm_model.save_model()
checkpoint = tf.train.Checkpoint(ugrnn_model1)
save_path = checkpoint.save(checkpoint_directory)

#####

f, (ax1, ax2) = plt.subplots(1, 2, sharey=True, figsize=(10, 4))
ax1.plot(range(len(lstm_model1.history['train_acc'])), lstm_model1.history['train_
acc'],
        label='LSTM Train Accuracy');
ax1.plot(range(len(lstm_model1.history['eval_acc'])), lstm_model1.history['eval_ac
c'],
        label='LSTM Test Accuracy');
ax2.plot(range(len(ugrnn_model1.history['train_acc'])), ugrnn_model1.history['trai
n_acc'],
```

```
        label='UGRNN Train Accuracy');  
ax2.plot(range(len(ugrnn_model1.history['eval_acc'])), ugrnn_model1.history['eval_'  
acc'],  
        label='UGRNN Test Accuracy');  
ax1.legend();  
ax2.legend();
```

WARNING:tensorflow:<keras.layers.legacy\_rnn.rnn\_cell\_impl.BasicLSTMCell object at 0x7fec2908f850>: Note that this cell is not optimized for performance. Please use tf.contrib.cudnn\_rnn.CudnnLSTM for better performance on GPU.

/tmp/ipykernel\_7766/2071740490.py:45: UserWarning: `tf.nn.rnn\_cell.BasicLSTMCell` is deprecated and will be removed in a future version. This class is equivalent as `tf.keras.layers.LSTMCell`, and will be replaced by that in Tensorflow 2.0.

```
self.rnn_cell = tf.compat.v1.nn.rnn_cell.BasicLSTMCell(cell_size)
/home/skanda/Softwares/miniconda3/lib/python3.9/site-packages/keras/layers/legacy_rnn/rnn_cell_impl.py:754: UserWarning: `layer.add_variable` is deprecated and will be removed in a future version. Please use `layer.add_weight` method instead.
```

```
self._kernel = self.add_variable(
/home/skanda/Softwares/miniconda3/lib/python3.9/site-packages/keras/layers/legacy_rnn/rnn_cell_impl.py:757: UserWarning: `layer.add_variable` is deprecated and will be removed in a future version. Please use `layer.add_weight` method instead.
```

```
self._bias = self.add_variable(
```

Train accuracy at epoch 1: 0.6242353  
Eval accuracy at epoch 1: 0.61653334  
Train accuracy at epoch 2: 0.73778826  
Eval accuracy at epoch 2: 0.7410667  
Train accuracy at epoch 3: 0.8689412  
Eval accuracy at epoch 3: 0.86693335  
Train accuracy at epoch 4: 0.89675295  
Eval accuracy at epoch 4: 0.8992  
Train accuracy at epoch 5: 0.9232941  
Eval accuracy at epoch 5: 0.92186666  
Train accuracy at epoch 6: 0.9458824  
Eval accuracy at epoch 6: 0.9461333  
Train accuracy at epoch 7: 0.9581647  
Eval accuracy at epoch 7: 0.9632  
Train accuracy at epoch 8: 0.9687059  
Eval accuracy at epoch 8: 0.976  
Train accuracy at epoch 9: 0.9760941  
Eval accuracy at epoch 9: 0.97653335  
Train accuracy at epoch 10: 0.9824  
Eval accuracy at epoch 10: 0.98373336  
Train accuracy at epoch 11: 0.9867294  
Eval accuracy at epoch 11: 0.9888  
Train accuracy at epoch 12: 0.9902118  
Eval accuracy at epoch 12: 0.9912  
Train accuracy at epoch 13: 0.9923294  
Eval accuracy at epoch 13: 0.99306667  
Train accuracy at epoch 14: 0.99524707  
Eval accuracy at epoch 14: 0.9949333  
Train accuracy at epoch 15: 0.99595296  
Eval accuracy at epoch 15: 0.99546665  
Train accuracy at epoch 16: 0.9960471  
Eval accuracy at epoch 16: 0.9970667  
Train accuracy at epoch 17: 0.9972706  
Eval accuracy at epoch 17: 0.99866664  
Train accuracy at epoch 18: 0.9979294  
Eval accuracy at epoch 18: 0.9997333  
Train accuracy at epoch 19: 0.9986824  
Eval accuracy at epoch 19: 0.99946666  
Train accuracy at epoch 20: 0.9987765  
Eval accuracy at epoch 20: 0.9997333  
Test accuracy: 0.8564

WARNING:tensorflow:<keras.layerslegacy\_rnn.rnn\_cell\_impl.BasicRNNCell object at 0x7febc0368520>: Note that this cell is not optimized for performance. Please use tf.contrib.cudnn\_rnn.CudnnRNNTanh for better performance on GPU.

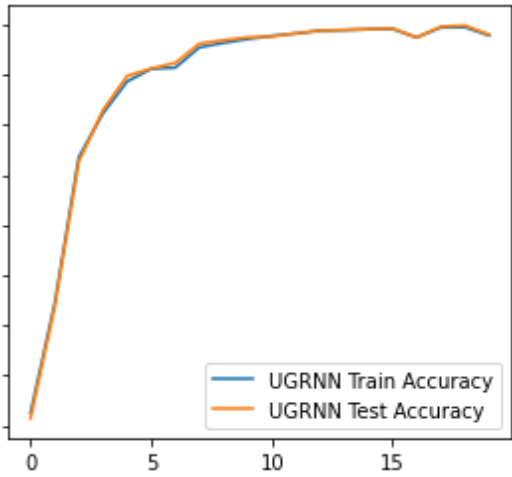
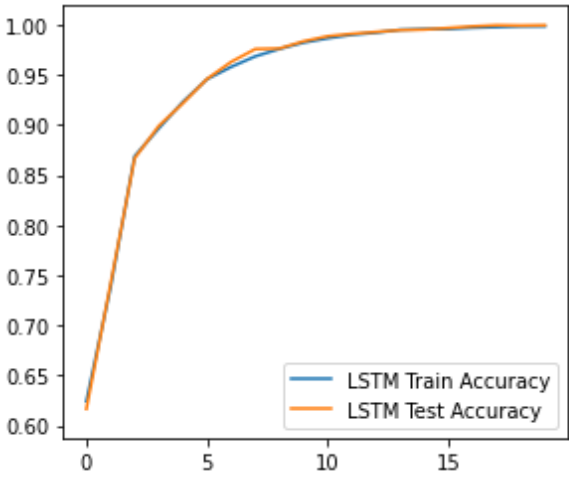
```
/tmp/ipykernel_7766/2071740490.py:48: UserWarning: `tf.nn.rnn_cell.BasicRNNCell` is deprecated and will be removed in a future version. This class is equivalent as `tf.keras.layers.SimpleRNNCell`, and will be replaced by that in Tensorflow 2.0.
```

```
self.rnn_cell = tf.compat.v1.nn.rnn_cell.BasicRNNCell(cell_size)
/home/skanda/Softwares/miniconda3/lib/python3.9/site-packages/keras/layers/legacy_rnn/rnn_cell_impl.py:457: UserWarning: `layer.add_variable` is deprecated and will be removed in a future version. Please use `layer.add_weight` method instead.
```

```
self._kernel = self.add_variable(
/home/skanda/Softwares/miniconda3/lib/python3.9/site-packages/keras/layers/legacy_rnn/rnn_cell_impl.py:460: UserWarning: `layer.add_variable` is deprecated and will be removed in a future version. Please use `layer.add_weight` method instead.
```

```
self._bias = self.add_variable(
```

Train accuracy at epoch 1: 0.61242354  
Eval accuracy at epoch 1: 0.6072  
Train accuracy at epoch 2: 0.72145885  
Eval accuracy at epoch 2: 0.7184  
Train accuracy at epoch 3: 0.86804706  
Eval accuracy at epoch 3: 0.86373335  
Train accuracy at epoch 4: 0.9117647  
Eval accuracy at epoch 4: 0.9149333  
Train accuracy at epoch 5: 0.94357646  
Eval accuracy at epoch 5: 0.94906664  
Train accuracy at epoch 6: 0.9562353  
Eval accuracy at epoch 6: 0.9565333  
Train accuracy at epoch 7: 0.95741177  
Eval accuracy at epoch 7: 0.96213335  
Train accuracy at epoch 8: 0.9776941  
Eval accuracy at epoch 8: 0.9813333  
Train accuracy at epoch 9: 0.98211765  
Eval accuracy at epoch 9: 0.9848  
Train accuracy at epoch 10: 0.98597646  
Eval accuracy at epoch 10: 0.9877333  
Train accuracy at epoch 11: 0.9891294  
Eval accuracy at epoch 11: 0.9885333  
Train accuracy at epoch 12: 0.9915294  
Eval accuracy at epoch 12: 0.9917333  
Train accuracy at epoch 13: 0.99435294  
Eval accuracy at epoch 13: 0.9944  
Train accuracy at epoch 14: 0.9948235  
Eval accuracy at epoch 14: 0.9952  
Train accuracy at epoch 15: 0.99576473  
Eval accuracy at epoch 15: 0.996  
Train accuracy at epoch 16: 0.9961412  
Eval accuracy at epoch 16: 0.99653333  
Train accuracy at epoch 17: 0.9875294  
Eval accuracy at epoch 17: 0.9874667  
Train accuracy at epoch 18: 0.99755293  
Eval accuracy at epoch 18: 0.9984  
Train accuracy at epoch 19: 0.9978353  
Eval accuracy at epoch 19: 0.99946666  
Train accuracy at epoch 20: 0.9895059  
Eval accuracy at epoch 20: 0.9904  
Test accuracy: 0.84372





In [8]:

```
# Specify the path where you want to save/restore the trained variables.
checkpoint_directory = 'models_checkpoints/ImdbRNN2/'

# Use the GPU if available.
device = 'gpu:0'

# Define optimizer.
optimizer = tf.compat.v1.train.AdamOptimizer(learning_rate=1e-5)

# Instantiate model. This doesn't initialize the variables yet.
lstm_model2 = RNNModel(vocabulary_size=len(word2idx), device=device,
                        checkpoint_directory=checkpoint_directory)

# Train model
lstm_model2.fit(train_dataset, val_dataset, test_dataset, optimizer, num_epochs=20,
                early_stopping_rounds=5, verbose=1, train_from_scratch=True)

#lstm_model.save_model()
checkpoint = tf.train.Checkpoint(lstm_model2)
save_path = checkpoint.save(checkpoint_directory)

#####
#####

# Define optimizer.
optimizer = tf.compat.v1.train.AdamOptimizer(learning_rate=1e-5)

# Instantiate model. This doesn't initialize the variables yet.
ugrnn_model2 = RNNModel(vocabulary_size=len(word2idx), rnn_cell='ugrnn',
                        device=device, checkpoint_directory=checkpoint_directory)

# Train model
ugrnn_model2.fit(train_dataset, val_dataset, test_dataset, optimizer, num_epochs=
20,
                early_stopping_rounds=5, verbose=1, train_from_scratch=True)

#lstm_model.save_model()
checkpoint = tf.train.Checkpoint(ugrnn_model2)
save_path = checkpoint.save(checkpoint_directory)

f, (ax1, ax2) = plt.subplots(1, 2, sharey=True, figsize=(10, 4))
ax1.plot(range(len(lstm_model2.history['train_acc'])), lstm_model2.history['train_
acc'],
        label='LSTM Train Accuracy');
ax1.plot(range(len(lstm_model2.history['eval_acc'])), lstm_model2.history['eval_ac
c'],
        label='LSTM Test Accuracy');
ax2.plot(range(len(ugrnn_model2.history['train_acc'])), ugrnn_model2.history['trai
n_acc'],
        label='UGRNN Train Accuracy');
ax2.plot(range(len(ugrnn_model2.history['eval_acc'])), ugrnn_model2.history['eval_
```

```
acc'],  
        label='UGRNN Test Accuracy');  
ax1.legend();  
ax2.legend();
```

WARNING:tensorflow:<keras.layers.legacy\_rnn.rnn\_cell\_impl.BasicLSTMCell object at 0x7fec291658e0>: Note that this cell is not optimized for performance. Please use tf.contrib.cudnn\_rnn.CudnnLSTM for better performance on GPU.

/tmp/ipykernel\_7766/2071740490.py:45: UserWarning: `tf.nn.rnn\_cell.BasicLSTMCell` is deprecated and will be removed in a future version. This class is equivalent to `tf.keras.layers.LSTMCell`, and will be replaced by that in Tensorflow 2.0.

```
self.rnn_cell = tf.compat.v1.nn.rnn_cell.BasicLSTMCell(cell_size)
```



Train accuracy at epoch 1: 0.5209412  
Eval accuracy at epoch 1: 0.52746665  
Train accuracy at epoch 2: 0.5607529  
Eval accuracy at epoch 2: 0.54826665  
Train accuracy at epoch 3: 0.5968  
Eval accuracy at epoch 3: 0.5944  
Train accuracy at epoch 4: 0.59115297  
Eval accuracy at epoch 4: 0.5832  
Train accuracy at epoch 5: 0.6383059  
Eval accuracy at epoch 5: 0.6242667  
Train accuracy at epoch 6: 0.6759059  
Eval accuracy at epoch 6: 0.68186665  
Train accuracy at epoch 7: 0.69534117  
Eval accuracy at epoch 7: 0.6821333  
Train accuracy at epoch 8: 0.7112  
Eval accuracy at epoch 8: 0.7088  
Train accuracy at epoch 9: 0.7168941  
Eval accuracy at epoch 9: 0.7154667  
Train accuracy at epoch 10: 0.7405647  
Eval accuracy at epoch 10: 0.7376  
Train accuracy at epoch 11: 0.76  
Eval accuracy at epoch 11: 0.75733334  
Train accuracy at epoch 12: 0.6872941  
Eval accuracy at epoch 12: 0.6810667  
Train accuracy at epoch 13: 0.82305884  
Eval accuracy at epoch 13: 0.82453334  
Train accuracy at epoch 14: 0.81590587  
Eval accuracy at epoch 14: 0.8189333  
Train accuracy at epoch 15: 0.83651763  
Eval accuracy at epoch 15: 0.84026664  
Train accuracy at epoch 16: 0.83331764  
Eval accuracy at epoch 16: 0.8509333  
Train accuracy at epoch 17: 0.84197646  
Eval accuracy at epoch 17: 0.848  
Train accuracy at epoch 18: 0.84381175  
Eval accuracy at epoch 18: 0.8498667  
Train accuracy at epoch 19: 0.84922355  
Eval accuracy at epoch 19: 0.85653335  
Train accuracy at epoch 20: 0.8481412  
Eval accuracy at epoch 20: 0.8549333  
Test accuracy: 0.7908

WARNING:tensorflow:<keras.layers.legacy\_rnn.rnn\_cell\_impl.BasicRNNCell object at 0x7febc03732b0>: Note that this cell is not optimized for performance. Please use tf.contrib.cudnn\_rnn.CudnnRNNTanh for better performance on GPU.

/tmp/ipykernel\_7766/2071740490.py:48: UserWarning: `tf.nn.rnn\_cell.BasicRNNCell` is deprecated and will be removed in a future version. This class is equivalent as `tf.keras.layers.SimpleRNNCell`, and will be replaced by that in Tensorflow 2.0.

```
self.rnn_cell = tf.compat.v1.nn.rnn_cell.BasicRNNCell(cell_size)
```

Train accuracy at epoch 1: 0.54992944  
Eval accuracy at epoch 1: 0.54826665  
Train accuracy at epoch 2: 0.5930353  
Eval accuracy at epoch 2: 0.5933333  
Train accuracy at epoch 3: 0.6171294  
Eval accuracy at epoch 3: 0.60746664  
Train accuracy at epoch 4: 0.67877644  
Eval accuracy at epoch 4: 0.6842667  
Train accuracy at epoch 5: 0.73505884  
Eval accuracy at epoch 5: 0.7389333  
Train accuracy at epoch 6: 0.7357647  
Eval accuracy at epoch 6: 0.73733336  
Train accuracy at epoch 7: 0.76950586  
Eval accuracy at epoch 7: 0.7744  
Train accuracy at epoch 8: 0.7637647  
Eval accuracy at epoch 8: 0.77466667  
Train accuracy at epoch 9: 0.78018826  
Eval accuracy at epoch 9: 0.7877333  
Train accuracy at epoch 10: 0.7934588  
Eval accuracy at epoch 10: 0.7962667  
Train accuracy at epoch 11: 0.80988234  
Eval accuracy at epoch 11: 0.8114667  
Train accuracy at epoch 12: 0.8231059  
Eval accuracy at epoch 12: 0.8354667  
Train accuracy at epoch 13: 0.8207059  
Eval accuracy at epoch 13: 0.82586664  
Train accuracy at epoch 14: 0.8296  
Eval accuracy at epoch 14: 0.8296  
Train accuracy at epoch 15: 0.83157647  
Eval accuracy at epoch 15: 0.8368  
Train accuracy at epoch 16: 0.8329412  
Eval accuracy at epoch 16: 0.83253336  
Train accuracy at epoch 17: 0.85237646  
Eval accuracy at epoch 17: 0.84933335  
Train accuracy at epoch 18: 0.8495059  
Eval accuracy at epoch 18: 0.85786664  
Train accuracy at epoch 19: 0.8512941  
Eval accuracy at epoch 19: 0.85866666  
Train accuracy at epoch 20: 0.8531294  
Eval accuracy at epoch 20: 0.8549333  
Test accuracy: 0.75392



In [9]:

```
# Specify the path where you want to save/restore the trained variables.
checkpoint_directory = './models_checkpoints/ImdbRNN3/'

# Use the GPU if available.
device = 'gpu:0'

# Define optimizer.
optimizer = tf.compat.v1.train.GradientDescentOptimizer(learning_rate=1e-3)

# Instantiate model. This doesn't initialize the variables yet.
lstm_model3 = RNNModel(vocabulary_size=len(word2idx), device=device,
                        checkpoint_directory=checkpoint_directory)

# Train model
lstm_model3.fit(train_dataset, val_dataset, test_dataset, optimizer, num_epochs=20,
                early_stopping_rounds=5, verbose=1, train_from_scratch=True)

#lstm_model.save_model()
checkpoint = tf.train.Checkpoint(lstm_model3)
save_path = checkpoint.save(checkpoint_directory)

#####

# Define optimizer.
optimizer = tf.compat.v1.train.GradientDescentOptimizer(learning_rate=1e-3)

# Instantiate model. This doesn't initialize the variables yet.
ugrnn_model3 = RNNModel(vocabulary_size=len(word2idx), rnn_cell='ugrnn',
                        device=device, checkpoint_directory=checkpoint_directory)

# Train model
ugrnn_model3.fit(train_dataset, val_dataset, test_dataset, optimizer, num_epochs=
20,
                early_stopping_rounds=5, verbose=1, train_from_scratch=True)

#lstm_model.save_model()
checkpoint = tf.train.Checkpoint(ugrnn_model3)
save_path = checkpoint.save(checkpoint_directory)

#####

f, (ax1, ax2) = plt.subplots(1, 2, sharey=True, figsize=(10, 4))
ax1.plot(range(len(lstm_model3.history['train_acc'])), lstm_model3.history['train_
acc'],
        label='LSTM Train Accuracy');
ax1.plot(range(len(lstm_model3.history['eval_acc'])), lstm_model3.history['eval_ac
c'],
        label='LSTM Test Accuracy');
ax2.plot(range(len(ugrnn_model3.history['train_acc'])), ugrnn_model3.history['trai
n_acc'],
```

```
        label='UGRNN Train Accuracy');  
ax2.plot(range(len(ugrnn_model3.history['eval_acc'])), ugrnn_model3.history['eval_'  
acc'],  
        label='UGRNN Test Accuracy');  
ax1.legend();  
ax2.legend();
```



WARNING:tensorflow:<keras.layers.legacy\_rnn.rnn\_cell\_impl.BasicLSTMCell object at 0x7febc0377bb0>: Note that this cell is not optimized for performance. Please use tf.contrib.cudnn\_rnn.CudnnLSTM for better performance on GPU.

/tmp/ipykernel\_7766/2071740490.py:45: UserWarning: `tf.nn.rnn\_cell.BasicLSTMCell` is deprecated and will be removed in a future version. This class is equivalent as `tf.keras.layers.LSTMCell`, and will be replaced by that in Tensorflow 2.0.

```
self.rnn_cell = tf.compat.v1.nn.rnn_cell.BasicLSTMCell(cell_size)
```

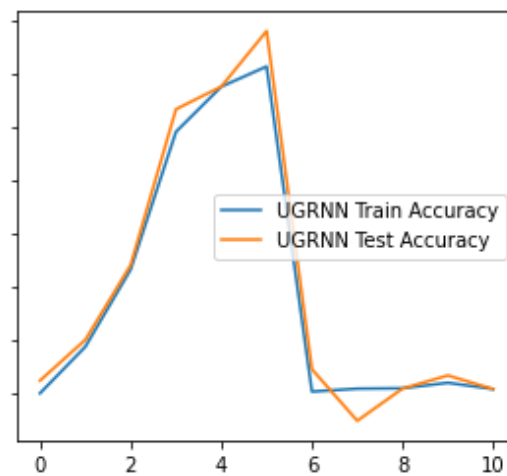

```
Train accuracy at epoch 1: 0.50037646
Eval accuracy at epoch 1: 0.49946666
Train accuracy at epoch 2: 0.5010353
Eval accuracy at epoch 2: 0.49333334
Train accuracy at epoch 3: 0.50032943
Eval accuracy at epoch 3: 0.49013335
Train accuracy at epoch 4: 0.50164706
Eval accuracy at epoch 4: 0.48853335
Train accuracy at epoch 5: 0.5000471
Eval accuracy at epoch 5: 0.5072
Train accuracy at epoch 6: 0.5000471
Eval accuracy at epoch 6: 0.4976
Train accuracy at epoch 7: 0.5209412
Eval accuracy at epoch 7: 0.5104
Train accuracy at epoch 8: 0.5130353
Eval accuracy at epoch 8: 0.51893336
Train accuracy at epoch 9: 0.50235295
Eval accuracy at epoch 9: 0.50346667
Train accuracy at epoch 10: 0.51981175
Eval accuracy at epoch 10: 0.5146667
Train accuracy at epoch 11: 0.59091765
Eval accuracy at epoch 11: 0.6002667
Train accuracy at epoch 12: 0.58174115
Eval accuracy at epoch 12: 0.59066665
Train accuracy at epoch 13: 0.6030588
Eval accuracy at epoch 13: 0.5976
Train accuracy at epoch 14: 0.54771763
Eval accuracy at epoch 14: 0.5477333
Train accuracy at epoch 15: 0.5787294
Eval accuracy at epoch 15: 0.58933336
Train accuracy at epoch 16: 0.55651766
Eval accuracy at epoch 16: 0.5648
Test accuracy: 0.5558
```

WARNING:tensorflow:<keras.layers.legacy\_rnn.rnn\_cell\_impl.BasicRNNCell object at 0x7fec291f6c10>: Note that this cell is not optimized for performance. Please use tf.contrib.cudnn\_rnn.CudnnRNNTanh for better performance on GPU.

/tmp/ipykernel\_7766/2071740490.py:48: UserWarning: `tf.nn.rnn\_cell.BasicRNNCell` is deprecated and will be removed in a future version. This class is equivalent as `tf.keras.layers.SimpleRNNCell`, and will be replaced by that in Tensorflow 2.0.

```
self.rnn_cell = tf.compat.v1.nn.rnn_cell.BasicRNNCell(cell_size)
```

Train accuracy at epoch 1: 0.49962354  
Eval accuracy at epoch 1: 0.5056  
Train accuracy at epoch 2: 0.5215529  
Eval accuracy at epoch 2: 0.5248  
Train accuracy at epoch 3: 0.5579294  
Eval accuracy at epoch 3: 0.56  
Train accuracy at epoch 4: 0.6227294  
Eval accuracy at epoch 4: 0.6333333  
Train accuracy at epoch 5: 0.64395297  
Eval accuracy at epoch 5: 0.644  
Train accuracy at epoch 6: 0.65345883  
Eval accuracy at epoch 6: 0.67013335  
Train accuracy at epoch 7: 0.50042355  
Eval accuracy at epoch 7: 0.51093334  
Train accuracy at epoch 8: 0.50178826  
Eval accuracy at epoch 8: 0.48666668  
Train accuracy at epoch 9: 0.5019765  
Eval accuracy at epoch 9: 0.50186664  
Train accuracy at epoch 10: 0.5045647  
Eval accuracy at epoch 10: 0.508  
Train accuracy at epoch 11: 0.5016  
Eval accuracy at epoch 11: 0.5016  
Test accuracy: 0.50148



In [19]:

```
# Specify the path where you want to save/restore the trained variables.
checkpoint_directory = './models_checkpoints/ImdbRNN4/'

# Use the GPU if available.
device = 'gpu:0'

# Define optimizer.
optimizer = tf.compat.v1.train.GradientDescentOptimizer(learning_rate=1e-4)

# Instantiate model. This doesn't initialize the variables yet.
lstm_model4 = RNNModel(vocabulary_size=len(word2idx), device=device,
                        checkpoint_directory=checkpoint_directory)

# Train model
lstm_model4.fit(train_dataset, val_dataset, test_dataset, optimizer, num_epochs=20,
                early_stopping_rounds=5, verbose=1, train_from_scratch=True)

#lstm_model.save_model()
checkpoint = tf.train.Checkpoint(lstm_model4)
save_path = checkpoint.save(checkpoint_directory)

#####

# Define optimizer.
optimizer = tf.compat.v1.train.GradientDescentOptimizer(learning_rate=1e-4)

# Instantiate model. This doesn't initialize the variables yet.
ugrnn_model4 = RNNModel(vocabulary_size=len(word2idx), rnn_cell='ugrnn',
                        device=device, checkpoint_directory=checkpoint_directory)

# Train model
ugrnn_model4.fit(train_dataset, val_dataset, test_dataset, optimizer, num_epochs=
20,
                early_stopping_rounds=5, verbose=1, train_from_scratch=True)

#lstm_model.save_model()
checkpoint = tf.train.Checkpoint(ugrnn_model4)
save_path = checkpoint.save(checkpoint_directory)

#####

f, (ax1, ax2) = plt.subplots(1, 2, sharey=True, figsize=(10, 4))
ax1.plot(range(len(lstm_model4.history['train_acc'])), lstm_model4.history['train_
acc'],
        label='LSTM Train Accuracy');
ax1.plot(range(len(lstm_model4.history['eval_acc'])), lstm_model4.history['eval_ac
c'],
        label='LSTM Test Accuracy');
ax2.plot(range(len(ugrnn_model4.history['train_acc'])), ugrnn_model4.history['trai
n_acc'],
```

```
        label='UGRNN Train Accuracy');  
ax2.plot(range(len(ugrnn_model4.history['eval_acc'])), ugrnn_model4.history['eval_'  
acc'],  
        label='UGRNN Test Accuracy');  
ax1.legend();  
ax2.legend();
```

WARNING:tensorflow:<keras.layers.legacy\_rnn.rnn\_cell\_impl.BasicLSTMCell object at 0x7fec291e5040>: Note that this cell is not optimized for performance. Please use tf.contrib.cudnn\_rnn.CudnnLSTM for better performance on GPU.

/tmp/ipykernel\_7766/2071740490.py:45: UserWarning: `tf.nn.rnn\_cell.BasicLSTMCell` is deprecated and will be removed in a future version. This class is equivalent to `tf.keras.layers.LSTMCell`, and will be replaced by that in Tensorflow 2.0.

```
self.rnn_cell = tf.compat.v1.nn.rnn_cell.BasicLSTMCell(cell_size)
```



```
Train accuracy at epoch 1: 0.5167059
Eval accuracy at epoch 1: 0.51493335
Train accuracy at epoch 2: 0.5180706
Eval accuracy at epoch 2: 0.5074667
Train accuracy at epoch 3: 0.49981177
Eval accuracy at epoch 3: 0.4992
Train accuracy at epoch 4: 0.5246588
Eval accuracy at epoch 4: 0.52026665
Train accuracy at epoch 5: 0.5005647
Eval accuracy at epoch 5: 0.49733335
Train accuracy at epoch 6: 0.55764705
Eval accuracy at epoch 6: 0.5477333
Train accuracy at epoch 7: 0.5660706
Eval accuracy at epoch 7: 0.5629333
Train accuracy at epoch 8: 0.57468235
Eval accuracy at epoch 8: 0.5784
Train accuracy at epoch 9: 0.5579765
Eval accuracy at epoch 9: 0.5576
Train accuracy at epoch 10: 0.5000941
Eval accuracy at epoch 10: 0.5072
Train accuracy at epoch 11: 0.5154353
Eval accuracy at epoch 11: 0.52026665
Train accuracy at epoch 12: 0.5060706
Eval accuracy at epoch 12: 0.5232
Train accuracy at epoch 13: 0.5807059
Eval accuracy at epoch 13: 0.59893334
Train accuracy at epoch 14: 0.5139294
Eval accuracy at epoch 14: 0.5136
Train accuracy at epoch 15: 0.5845647
Eval accuracy at epoch 15: 0.58426666
Train accuracy at epoch 16: 0.51261175
Eval accuracy at epoch 16: 0.5104
Train accuracy at epoch 17: 0.54545885
Eval accuracy at epoch 17: 0.5498667
Train accuracy at epoch 18: 0.49957648
Eval accuracy at epoch 18: 0.49653333
Test accuracy: 0.5022
```

WARNING:tensorflow:<keras.layers.legacy\_rnn.rnn\_cell\_impl.BasicRNNCell object at 0x7fec29165a30>: Note that this cell is not optimized for performance. Please use tf.contrib.cudnn\_rnn.CudnnRNNTanh for better performance on GPU.

```
/tmp/ipykernel_7766/2071740490.py:48: UserWarning: `tf.nn.rnn_cell.BasicRNNCell` is deprecated and will be removed in a future version. This class is equivalent as `tf.keras.layers.SimpleRNNCell`, and will be replaced by that in Tensorflow 2.0.
```

```
self.rnn_cell = tf.compat.v1.nn.rnn_cell.BasicRNNCell(cell_size)
```

```
Train accuracy at epoch 1: 0.52254117
Eval accuracy at epoch 1: 0.5256
Train accuracy at epoch 2: 0.5192
Eval accuracy at epoch 2: 0.52613336
Train accuracy at epoch 3: 0.52117646
Eval accuracy at epoch 3: 0.51093334
Train accuracy at epoch 4: 0.5349647
Eval accuracy at epoch 4: 0.5432
Train accuracy at epoch 5: 0.5258353
Eval accuracy at epoch 5: 0.5218667
Train accuracy at epoch 6: 0.5516235
Eval accuracy at epoch 6: 0.5536
Train accuracy at epoch 7: 0.53157645
Eval accuracy at epoch 7: 0.5498667
Train accuracy at epoch 8: 0.5423529
Eval accuracy at epoch 8: 0.5362667
Train accuracy at epoch 9: 0.5584
Eval accuracy at epoch 9: 0.5538667
Train accuracy at epoch 10: 0.58894116
Eval accuracy at epoch 10: 0.5808
Train accuracy at epoch 11: 0.53548235
Eval accuracy at epoch 11: 0.54106665
Train accuracy at epoch 12: 0.6243765
Eval accuracy at epoch 12: 0.63706666
Train accuracy at epoch 13: 0.61331767
Eval accuracy at epoch 13: 0.61013335
Train accuracy at epoch 14: 0.6248
Eval accuracy at epoch 14: 0.62186664
Train accuracy at epoch 15: 0.6062588
Eval accuracy at epoch 15: 0.6016
Train accuracy at epoch 16: 0.6135059
Eval accuracy at epoch 16: 0.6138667
Train accuracy at epoch 17: 0.6120941
Eval accuracy at epoch 17: 0.6088
Test accuracy: 0.59728
```



In [ ]:

In [ ]:

In [ ]:

In [11]:

```
#####
# Import/download necessary libraries to process new sequences
#####
import nltk
try:
    nltk.data.find('tokenizers/punkt')
except LookupError:
    nltk.download('punkt')
from nltk.tokenize import word_tokenize
import re
```

In [12]:

```
def process_new_review(review):
    '''Function to process a new review.
    Args:
        review: original text review, string.
    Returns:
        indexed_review: sequence of integers, words correspondence
                        from word2idx.
        seq_length: the length of the review.
    '''
    indexed_review = re.sub(r'<[^>]+>', ' ', review)
    indexed_review = word_tokenize(indexed_review)
    indexed_review = [word2idx[word] if word in list(word2idx.keys()) else
                      word2idx['Unknown_token'] for word in indexed_review]
    indexed_review = indexed_review + [word2idx['End_token']]
    seq_length = len(indexed_review)
    return indexed_review, seq_length
```

In [13]:

```
sent_dict = {0: 'negative', 1: 'positive'}
```

In [14]:

```
review_score_10 = "I think Bad Apples is a great time and I recommend! I enjoyed the opening, which gave way for the rest of the movie to occur. The main couple was very likable and I believed all of their interactions. They had great onscreen chemistry and made me laugh quite a few times! Keeping the girls in the masks but seeing them in action was something I loved. It kept a mystery to them throughout. I think the dialogue was great. The kills were fun. And the special surprise gore effect at the end was AWESOME!! I won't spoil that part ;) I also enjoyed how the movie wrapped up. It gave a very urban legends type feel of \"did you ever hear the story...\". Plus it leaves the door open for another film which I wouldn't mind at all. Long story short, I think if you take the film for what it is; a fun little horror flick, then you won't be disappointed! HaPpY eArLy HaLLoWeEn!"
```

In [15]:

```
review_score_4 = "A young couple comes to a small town, where the husband gets a job working in a hospital. The wife which you instantly hate or dislike works home, at the same time a horrible murders takes place in this small town by two masked killers. Bad Apples is just your typical B-horror movie with average acting (I give them that. Although you may get the idea that some of the actors are crazy-conervative Christians), but the script is just bad, and that's what destroys the film."
```



In [16]:

```
review_score_1 = "When you first start watching this movie, you can tell its going to be a painful ride. the audio is poor...the attacks by the \"girls\" are like going back in time, to watching the old rocky films, were blows never touched. the editing is poor with it aswell, example the actress in is the bath when her husband comes home, clearly you see her wearing a flesh coloured bra in the bath. no hints or spoilers, just wait till you find it in a bargain basket of cheap dvds in a couple of weeks"
```

In [17]:

```
new_reviews = [review_score_10, review_score_4, review_score_1]
scores = [10, 4, 1]
```

In [18]:

```
with tf.device(device):
    for original_review, score in zip(new_reviews, scores):
        indexed_review, seq_length = process_new_review(original_review)
        indexed_review = tf.reshape(tf.constant(indexed_review), (1,-1))
        seq_length = tf.reshape(tf.constant(seq_length), (1,))
        logits = lstm_model.predict(indexed_review, seq_length, False)
        pred = tf.argmax(logits, axis=1).numpy()[0]
        print('The sentiment for the review with score %d was found to be %s'
              %(score, sent_dict[pred]))
```

```
-----
-----
NameError                                Traceback (most recent call
last)
/tmp/ipykernel_7766/413355369.py in <module>
      4         indexed_review = tf.reshape(tf.constant(indexed_review
), (1,-1))
      5         seq_length = tf.reshape(tf.constant(seq_length), (1,))
----> 6         logits = lstm_model.predict(indexed_review, seq_length
, False)
      7         pred = tf.argmax(logits, axis=1).numpy()[0]
      8         print('The sentiment for the review with score %d was
found to be %s'
```

NameError: name 'lstm\_model' is not defined